

Drink Recommender

Architecture and design decisions - SplitEasy

September 2026

What it does

You give it four things - which state you are in, who you are drinking with, how many people and what budget - and it answers the question you actually have: which bottles do I walk out of the shop with, and what will that cost each of us.

The one design decision that matters

There is no language model anywhere in this feature.

That was a deliberate choice, not a limitation. A model asked to recommend drinks will happily invent a price for a brand in a state it has never seen, and the answer will read exactly like a real one. For a feature whose entire job is money, that failure is worse than useless - it is confidently wrong.

So every number here is either looked up in a table with a citation, or computed in code from your own recorded spending. Prices are read, arithmetic is done in Python, and the ranking is a sort you can read and argue with. Every figure on screen can be traced to a source URL or one of your own expenses.

The two data sources

The recommender stands on two legs, and they do different jobs.

- 1. Price table** `backend/liquor_prices.py` - a hand-curated table scraped from public state price listings. Answers 'what does this bottle cost here'.
- 2. Your history** The expenses already in the app. Answers 'what do these particular people actually drink, and what do they normally spend'.

The second is the one a generic app cannot copy. Knowing that you and Anubhav have had 18 sessions averaging Rs 1,458 and keep buying Old Monk is worth more than any amount of general knowledge about whisky.

Request flow



Step 4 - choosing the bottles

Spirits in India are sold in three sizes and everybody names them the same way. The recommender only ever suggests these:

180 ml	quarter
375 ml	half
750 ml	full

You pick how much the group is drinking between them, in those sizes - a total for the room, not an allowance per head. 180 ml across two people is 90 ml each, and the cards say so. Given that and a budget range, the recommender lists every bottle of that size priced inside the range, dearest first.

'Best' is approximated by price. Within one state and one size, a dearer bottle is the better one often enough to be useful, and price is the only quality signal in the data. So the list starts with the most you can afford and works down, rather than the cheapest thing that technically fits. A brand you already drink wins a tie.

Budget is a range rather than a ceiling, because 'around 500' is what people mean. It has to be at least Rs 60 wide: a price list that moves in fifties will not reliably contain anything inside a Rs 30 window, and an empty result from too narrow a range looks like a missing brand rather than a bad question. When a range does come back empty, the answer carries what that size actually costs in that state - 'they run Rs 310-940 here' - so a dead end tells you how to fix it.

Gurugram, quarter (180 ml) between 2 people, Rs 100-250

Antiquity Blue	Rs 250	90 ml each
McDowell's Single Malt	Rs 250	90 ml each
Smirnoff	Rs 200	90 ml each

An empty list distinguishes its two causes. Nothing published at that size in that state reads differently from everything at that size being over budget, because they need different things from you.

Strength, and what each person actually gets

Every card carries the drink's strength as % ABV, and every card says what the basket works out to per person - both in millilitres and in millilitres of pure alcohol. That last figure is the only fair way to compare a strong beer against a mild one, or beer against spirits.

Splitting the bottle across the group is shown too - 750 ml between four is 188 ml each - so the per-head cost and the per-head measure sit next to each other.

ABV is a property of the brand rather than the state, so unlike price it is a single table. Brands whose strength was actually published carry that figure; everything else falls back to the category typical - 42.8% for Indian-made spirits by excise convention, 40% for imported - and is prefixed with a tilde on screen so a typical is never presented as a label reading.

Beer

Beer is the fourth option in the same control as the three bottle sizes, because it is a fourth answer to the same question. Choosing it returns beer instead of spirits - showing beer alongside every spirit answer would have made choosing it mean nothing.

Beer cards are sized to the budget: two each is the default round, and if that overruns the budget the count comes down rather than the beer disappearing. They are ordered strongest first, so the list reads as a real choice between a session beer and a strong one.

Step 5 - the NCR comparison

Alcohol is a State subject in India. Every state sets its own excise duty and its own MRP, so the same bottle genuinely costs different amounts a short drive apart. That is not a rounding difference - it is often hundreds of rupees.

Royal Stag 750 ml	UP	650-720	Delhi	550-800	Gurugram	480-550
Blenders Pride 750 ml	UP	680-780	Delhi	550-800	Gurugram	730
Old Monk 750 ml	UP	520	Delhi	350-420	Gurugram	-

So every suggestion is priced in all three NCR regions and the cheapest is highlighted. A dash means that region publishes no price for that exact bottle and size. It is never a number carried across from a neighbouring state, even though that would fill the grid more neatly.

A region is only priced when it has every size in the basket. Pricing half a basket and estimating the rest would produce a total that reads as real and is not.

Honesty rules in the price table

- Every row cites a source URL and the month it was read.
- Nothing is interpolated between states. A missing state is a dash, not a guess.
- Where a source published a range, the range is kept rather than averaged away.
- Haryana sets a minimum selling price, not a fixed MRP, so shops legally differ and two honest sources disagreed. Those rows span both figures instead of one being picked and presented as the price.

Current coverage: Delhi, Gurugram (Haryana), Uttar Pradesh and Maharashtra. Running 'python liquor_prices.py' prints coverage and every source, so the gaps are visible rather than implied.

Privacy

History is computed only from groups the caller is a member of. Naming somebody you do not share a group with returns nothing about them - the endpoint cannot be used to read another person's drinking habits, in the same way the rest of the API cannot be used to read their balances.

Ranking

The sort is three keys, in this order:

- 1. Familiar first** Brands this set of people has actually bought before.
- 2. Right amount** Closest total volume to what this many people would get through.
- 3. Cheapest** Then price, ascending.

History only appears on screen once you have named somebody. With nobody named it was the caller's own drinking across every group, shown as 'sessions together' with no one to have had them with. The numbers still rank the suggestions in that case - they just are not presented as shared history.

Files

backend/liquor_prices.py	The price table, its sources, and the cross-region lookups. Pure data plus small helpers - no database, no framework.
backend/routers/recommend.py	Two endpoints. GET /recommend/meta lists the states we have prices for; GET /recommend/ does the work. Holds the history scan, the bottle solver and the ranking.
frontend/src/pages/Recommend.js x	The form and the result cards, including the three-region price strip.
frontend/src/components/BottomNav.jsx	The raised centre key routes here.

Known limits

- Uttar Pradesh coverage is thin - 8 brands - because UP publishes PDFs rather than anything machine-readable. Real shop prices would fix this fastest.
- Prices drift with every excise year, and shops vary within a state. These are indicative, which is what the source dates are for.
- Brand matching across regions is on name and size. A brand written differently by two sources will not match, and shows as a dash rather than a wrong price.
- Smaller bottles are thin outside Haryana: quarts are what most sources publish. Picking quarter or half in a state without them says so plainly.

What it would take to add a state

Add rows to liquor_prices.py with a real source, and it appears in the dropdown automatically. Nothing else changes: the solver, the comparison and the ranking are all driven off the table. If a state should join the NCR comparison, add it to the NCR tuple.